

Zero2Hero: Responsible AI

The honest AI cofounder for ambitious students. Responsibility is layered across input, reasoning, output, control, and privacy, and every layer is verified in the codebase, not just claimed.

- **Honesty is engineered, not prompted**

The strongest form of responsible AI here is architectural. The verdict can only say "Build it" when a high-risk assumption has actually passed and is backed by real-world field evidence, so it is gated in deterministic code and cannot be talked into a yes. A failed high-risk assumption forces a kill or pivot. (lib/verdict.ts)

- **Input safety**

Clearly harmful or illegal ideas (weapons, illegal drugs, malware, fraud, exploitation, violence) are refused server-side, before any model call, plus a prompt-level refusal. The screen is deliberately conservative to avoid false positives on legitimate ideas such as a "Nerf gun marketplace" or a "malware-detection tool." (lib/safety.ts)

- **Anti-sycophancy by design**

Zero2Hero is not a cheerleader. The interview hunts the weakest assumption, an adversarial cofounder argues against the idea, and a field-test judge scores results only against their own pre-registered pass/fail thresholds. (lib/prompts.ts, lib/verdict.ts)

- **Grounding and anti-hallucination**

The model is instructed never to invent facts, research uses only what the findings support, and fabricated citations are dropped server-side so a fake source can never move the confidence score. Only real URLs survive. (lib/research.ts)

- **Output integrity**

Every model response is validated against a Zod schema at the API boundary, returning a 422 on mismatch, so malformed output never reaches the user. This is what makes Zero2Hero a system rather than a chat wrapper. (lib/schema.ts)

- **Human-in-the-loop control**

The AI only ever proposes. The user applies every change (evidence, status, concede) and makes the final go / no-go call. Nothing auto-commits. It is a copilot for judgment, not an autopilot.

- **Privacy by default**

Runs fully local on Ollama with no data leaving the machine, uses a self-hosted SearxNG search with no third-party tracking, keeps keys server-side only (never in the client bundle), stores plans in browser localStorage with no server database, and rate-limits its routes.

- **Honest about limits**

A persistent disclaimer states that this is AI guidance, not professional, legal, or financial advice, and that users should validate independently.

Built for the USAI Global AI Hackathon 2026 (College track, AI for Life & Work) and its Responsible AI Award.

Zero2Hero. Local-first, evidence-driven, human-in-the-loop.